

ML Week 0x03 Feature Extraction

Categorical variables

- *One of K* or *one-hot* encoding — one binary feature per possible value
- The importance of not encoding order where none exists
- Text as explanatory variable \implies encode as feature vectors
- **Bag of words**
 - Corpus (collection of documents)
 - Vocabulary (set of unique words in document)
 - Words = dimensions
 - Order of words doesn't matter
 - Order in vectors encodes words
 - Binary: present or not
 - **CountVectorizer**
 - * by default converts to lowercase
 - * tokens
 - * stop words (mot vide) — words in most documents don't convey much information
 - * stemming — rule-based, drop suffixes
(racinisation ou désuffixation : transformer des flexions en leur radical or racine)
 - * lemmatization – find root form of word
(lemmatisation : transformer en lemme (forme canonique))
- **TF - IDF**
 - norm: L1, L2, none (in equation: max)
 - use_idf: enable IDF, default=True
 - smooth_idf: use $n_t + 1$ in IDF, default=True
 - sublinear_tf: Apply sublinear scaling, replacing TF_{td} with $1 + \log(TF_{td})$.
- **Exercise**: HashingVectorizer
 - Problems:
 1. Two passes to create structure: learn vocabulary (tokens), then create feature vectors

2. Vocabulary (dict) stored in memory

– Instead, `HashingVectorizer`

- * Bounded memory (no dict), even low memory (sparse scipy matrix)
- * Stateless, so can be used online (streaming) and parallel
- * Fast to serialize/unserialize
- * `n_features` defaults to 2^{20}
- * Note negative values. Increment takes sign of hash value, so possibility of cancellation.
- * **But** can't compute inverse transform, so hard to know which features are most important
- * Collisions can happen, but rare for 2^{20}
- * No IDF weighting, since IDF is stateful

OCR

– `sklearn.digits`

- * over 1700 hand-written digits (0–9)
- * 8×8 four bit pixels
- * white is most intense and represented by 0
- * black is least intense and represented by 16

– Feature vectors

- * In general, matrices not sparse
- * $100 \times 100 \implies 1e4$
- * $1920 \times 1080 \implies 2e6$
- * Problems: space, time
- * More problems: sensitive to position, rotation, scale
- * Even more problems: sensitive to illumination
- * We'll come back to this with SVM (machine à vecteurs de support)...

– Corners and edges

- * Basic computer vision techniques
- * Define: feature extraction
- * Define: feature engineering
- * Compression
- * Point matching
- * Edge detectors are mostly rotation invariant
- * So therefore corner detectors are, too
- * But scaling can hide corners

- SIFT
 - * Uses scale-space
 - * Approximates Laplacian of Gaussian with Difference of Gaussian for finding scale space
 - * Maybe a bit slower than we'd like
- SURF
 - * Speeds up SIFT
 - * Approximates Laplacian of Gaussian with Box Filter for finding scale space
 - * **Example**
 - * Standardized dataset: zero mean, unit variance (why?)
 - * $(x - \mu)/\sigma$