

ML Week 0x03 Logistic Regression

linear vs logistic ($\times 2$)

- y is discrete: classification
- Examples:
 - spam/non-spam
 - transaction: fraud or legitimate
 - tumor: malignant or benign
- So 0 or 1
- Problems with linear regression here (picture)
- sigmoid ($\times 2$)

Non-linear decision boundaries ($\times 2$)

- Still just use gradient descent
- This is why we like things to be differentiable
- Multinomial (multi-class) classification
 - one vs all (OvO, OvR) (draw picture, get three classifiers)
 - * At decision time, try $k - 1$ classifiers, choose the one with the most +1 votes (highest probability)
 - * Problem: learners see more negatives than positives
 - * Problem: different confidence for different decision boundaries
 - one vs one (OvO) (draw picture, get three classifiers)
 - * At decision time, try $k(k - 1)/2$ classifiers, choose the one with the most +1 votes (highest probability)

Cost function ($\times 7$)

- This is not convex
- So potentially many local minima

- Plot cost and explain what it means for $y \in \{0, 1\}$.
- Note that our convex cost function
 - is differentiable
 - can be derived from statistics using the principles of maximum likelihood estimation.